

Submitted by: Amadeus Hovekamp,
amadeus.hovekamp@rwth-aachen.de,
Matr. no: 4603934

The repository with the source code can be found under the following link:
<https://github.com/Schokokugel/RoboticsLab>

1 Running the default network and testing smaller batch sizes

The default settings for training the network are 12 training epochs using the complete training set with 50000 samples, a learning rate of 10^{-3} and batches of size 128. The two 2d convolution layers both consist of 16 filters of size 3×3 . The validation set contains 10000 and the final test set 10000 samples.

When I ran this network on my CPU for the first time, I achieved a very good performance even after only a single epoch, reaching more than 98 % accuracy. But when I managed to run the network on my GPU (NVIDIA GeForce GTX 950M, Ubuntu 18.04), the accuracy on the validation set reached barely 82 % after one epoch using the exact same code. Before I continued with task 2 I therefore checked, if the batch size of 128 might be the reason of low performance, because the resulting tensorflow graph might be too big to fit in the available 3.6 GB of graphic memory.

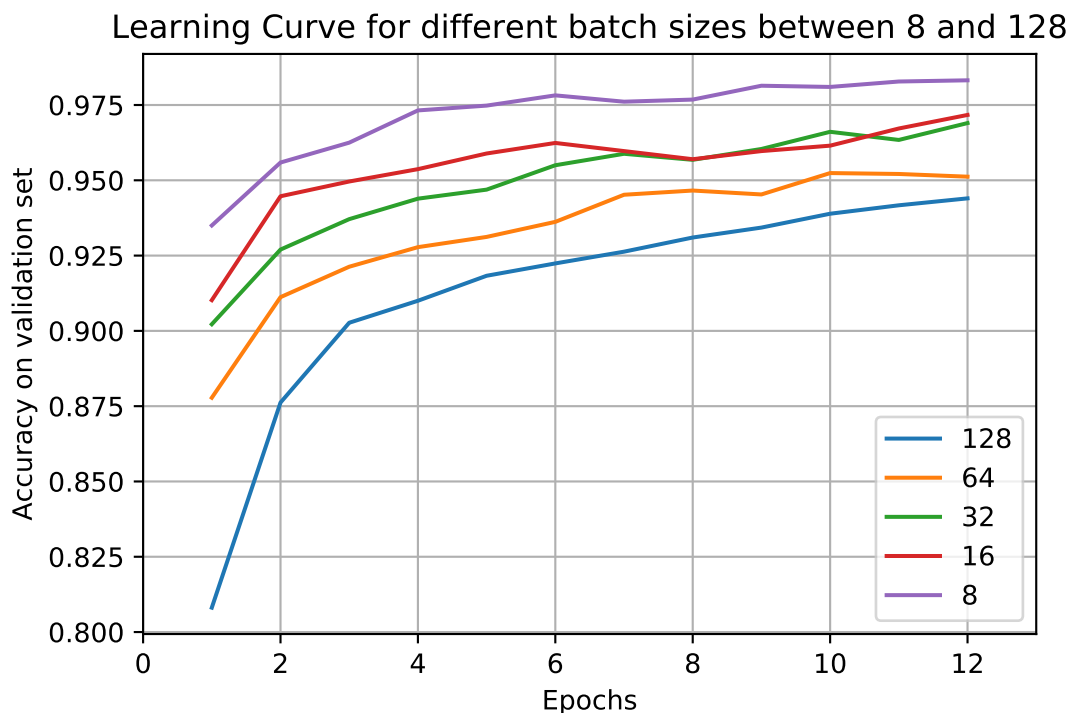


Figure 1: Learning curve for the default network

Figure 1 shows a connection between convergence speed and batch size, but where this difference comes from, I do not know. Using smaller batch sizes in the current implementation leads to more gradient steps in one epoch, which might be a reason for this performance difference.

2 Learning rate experiments

Even though the batch size had a big influence of the network's performance, I used the default settings for the experiments and only changed the learning rate. The results are shown in Figure 2.




Figure 2: Error Curve for different Learning Rates between 0.1 and 0.0001

One instantly sees the big difference between the smallest learning rate of 10^{-4} and the bigger learning rates. A learning rate that small might slow down the learning of a network very much. After one epoch the network barely performs better than random (which would yield an error rate about 0.9). After 12 epochs it still performs much worse than the other three networks after a single epoch of training. The other three networks perform similar but still one can see a trend: A bigger learning rate yields better results. Generally this is not true, because when the learning rate is set too big, the steps of the gradient descent algorithm might jump over valleys of the loss function and thus oscillate around a minimum or possibly even diverge if the step size is way too big.

3 Filter size experiments

In this experiment I only changed the filter size of the convolution layers. All filters in the two 2d convolution layers is of size $filtersize \times filtersize$. Figure 2 shows the results. The learning process of all networks looks similar but from the second epoch onwards one sees that a higher filter size leads to performance improvements.

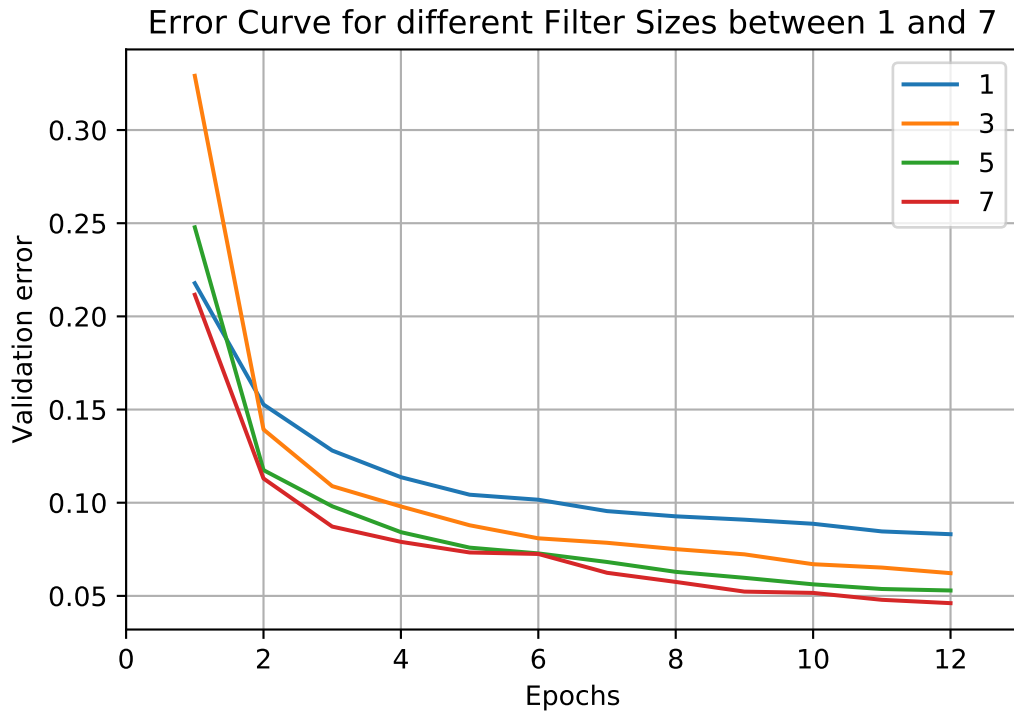


Figure 3: Error Curve for different Filter Sizes between 1 and 7

The higher filter size enables the network to recognize features of bigger size. A network with bigger filters is generally more powerful but also has more parameters which requires more computation to train it. Therefore the size of the filters used should match the size of the features those filters have to detect: Small features \rightarrow small filters. Large features \rightarrow large filters.

Deep convolution networks might learn these bigger features by combining several smaller features from the layer before. This way the filters of size 3 and 5 perform just a little worse than the filter of size 7.

4 Random search

In Figure 4 the results of 50 iterations of random search are shown. The red line displays the performance of the best configuration trained for 6 epochs found upto that iteration. Some somewhat good networks have been found by the random search, the best iteration being iteration 18.

When comparing the results of the best iterations of random search with the networks tested in task 1 and task 2 one can see that only the default network with a learning rate of 0.1 achieved similarly good performances on the validation set. This shows that when there is little knowledge of how the hyperparameters should be set for a specific problem, random search is a very useful approach to find good hyperparameter settings. One big downside is the long computation time which in my case was 220 minutes for 50 iterations and 6 epochs each, about 4.4 minutes per iteration. One can notice too, that the last 32 iterations, more than half of the total number of iterations did not improve the best configuration found so far.

Figure 5 shows the learning curve of a network trained with the best set of parameters found in iteration 18 of the random search. It has a rather small filter size but a relatively big learning rate and medium sized batch size. The performance is nearly the same as in the random search iteration, but one can see oscillations in the validation error, probably due to the rather large learning rate. The error on the test set is 0.009599983 which is just as good as the performance on the validation set.




Figure 4: Validation Error for all 50 Iterations of Random Search

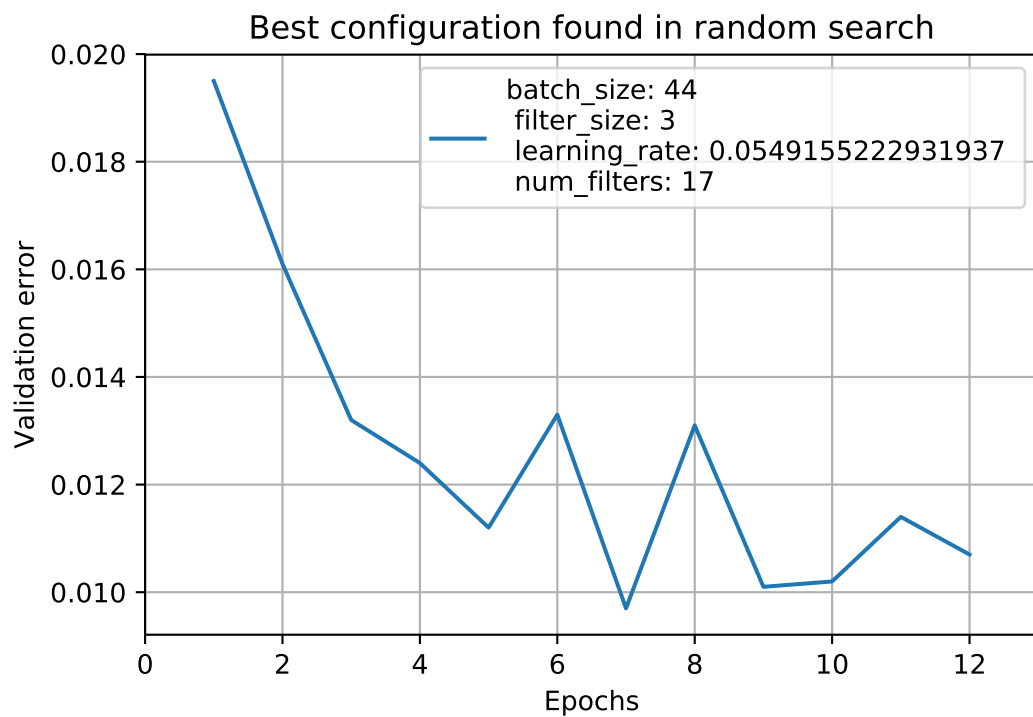


Figure 5: Validation Error of the best hyperparameter set found with Random Search